

Bank Loan Approval Prediction using Artificial Neural Network

In this project, we will build and train a deep neural network model to predict the likelihood of a liability customer buying personal loans based on customer features.

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.layers import Dense, Activation, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.metrics import Accuracy
import matplotlib.pyplot as plt

bank_df = pd.read_csv("UniversalBank.csv")

bank_df.head()
```

	ID	Age	Experience	Income	ZIP Code	Family	CCAvg	Education
0	1	25	1	49	91107	4	1.6	1
1	2	45	19	34	90089	3	1.5	1
2	3	39	15	11	94720	1	1.0	1
3	4	35	9	100	94112	1	2.7	2
4	5	35	8	45	91330	4	1.0	2

	Personal Loan	Securities Account	CD Account	Online	CreditCard
0	0	1	0	0	0
1	0	1	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	1

```
bank_df.shape

(5000, 14)
```

Exploratory Data Analysis

```
bank_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 14 columns):
-----
0    ID                5000 non-null    int64
1    Age               5000 non-null    int64
2    Experience         5000 non-null    int64
3    Income             5000 non-null    int64
4    ZIP Code           5000 non-null    int64
5    Family             5000 non-null    int64
6    CCAvg              5000 non-null    float64
7    Education          5000 non-null    int64
8    Mortgage           5000 non-null    int64
9    Personal Loan       5000 non-null    int64
10   Securities Account  5000 non-null    int64
11   CD Account         5000 non-null    int64
12   Online              5000 non-null    int64
13   CreditCard          5000 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 547.0 KB

bank_df.describe().transpose()

25% \
count      mean      std      min
ID      5000.0    2500.500000    1443.520003    1.0
Age      5000.0     45.338400     11.463166    23.0
```

Experience	5000.0	20.104600	11.467954	-3.0
10.00				
Income	5000.0	73.774200	46.033729	8.0
39.00				
ZIP Code	5000.0	93152.503000	2121.852197	9307.0
91911.00				
Family	5000.0	2.396400	1.147663	1.0
1.00				
CCAvg	5000.0	1.937938	1.747659	0.0
0.70				
Education	5000.0	1.881000	0.839869	1.0
1.00				
Mortgage	5000.0	56.498800	101.713802	0.0
0.00				
Personal Loan	5000.0	0.096000	0.294621	0.0
0.00				
Securities Account	5000.0	0.104400	0.305809	0.0
0.00				
CD Account	5000.0	0.060400	0.238250	0.0
0.00				
Online	5000.0	0.596800	0.490589	0.0
0.00				
CreditCard	5000.0	0.294000	0.455637	0.0
0.00				

	50%	75%	max
ID	2500.5	3750.25	5000.0
Age	45.0	55.00	67.0
Experience	20.0	30.00	43.0
Income	64.0	98.00	224.0
ZIP Code	93437.0	94608.00	96651.0
Family	2.0	3.00	4.0
CCAvg	1.5	2.50	10.0
Education	2.0	3.00	3.0
Mortgage	0.0	101.00	635.0
Personal Loan	0.0	0.00	1.0
Securities Account	0.0	0.00	1.0
CD Account	0.0	0.00	1.0
Online	1.0	1.00	1.0
CreditCard	0.0	1.00	1.0

```
bank_df.isnull().sum()
```

```

Age
Income
ZIP Code
Family
CCAvg
Education
Mortgage
Personal Loan
Securities Account
CD Account
Online
CreditCard
0

```

```
Education          0
Mortgage           0
Personal Loan      0
Securities Account 0
CD Account         0
Online            0
CreditCard        0
dtype: int64
```

```
avg_age = bank_df["Age"].mean()
print ("The average age of this dataset is {:.1f}.".format(avg_age))
```

The average age of this dataset is 45.3.

```
percent_cc = sum(bank_df["CreditCard"] == 1)/len(bank_df)
print ("The percentage of customers that own the bank's credit card is
{:.2%}.".format(percent_cc))
```

The percentage of customers that own the bank's credit card is 29.40%.

```
percent_loan = sum(bank_df["Personal Loan"] == 1)/len(bank_df)
print ("The percentage of customers that took out a personal loan is
{:.2%}.".format(percent_loan))
```

The percentage of customers that took out a personal loan is 9.60%.



Data Visualization


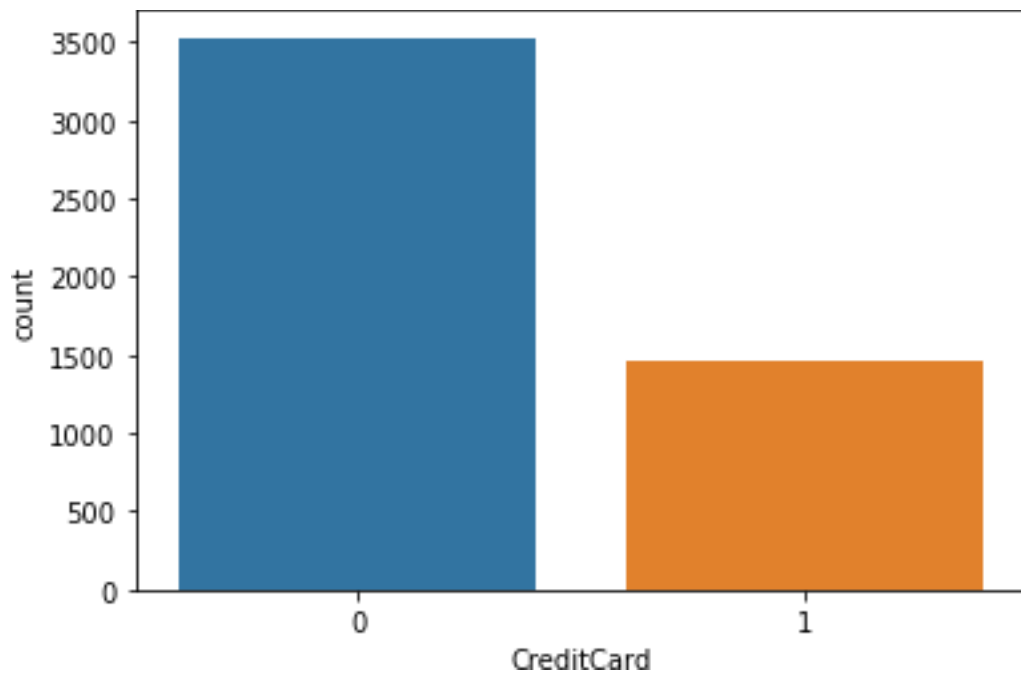
```
sns.countplot(x=bank_df["Personal Loan"])
plt.show()
```

```
sns.countplot(x=bank_df["Education"])
plt.show()
```


```
sns.countplot(x=bank_df["CreditCard"])
plt.show()
```

```
plt.figure(figsize=(20,10))
sns.countplot(x=bank_df["Age"])
plt.savefig('age.png', facecolor='w', bbox_inches='tight')
plt.show()
```





```
# lets look at the distribution of the income
plt.figure(figsize=(15,8))
sns.distplot(bank_df["Income"])
plt.savefig('income.png', facecolor='w', bbox_inches='tight')
plt.show()
```



```
personal_loans = bank_df[bank_df['Personal Loan'] == 1].copy()
no_personal_loans = bank_df[bank_df['Personal Loan'] == 0].copy()
```

```
personal_loans.describe().T
```

	count	mean	std	min
25% \				
ID	480.0	2390.650000	1394.393674	10.0
1166.50				
Age	480.0	45.066667	11.590964	26.0
35.00				
Experience	480.0	19.843750	11.582443	0.0
9.00				
Income	480.0	144.745833	31.584429	60.0
122.00				
ZIP Code	480.0	93153.202083	1759.223753	90016.0
91908.75				
Family	480.0	2.612500	1.115393	1.0
2.00				
CCAvg	480.0	3.905354	2.097681	0.0
2.60				
Education	480.0	2.233333	0.753373	1.0
2.00				
Mortgage	480.0	100.845833	160.847862	0.0
0.00				
Personal Loan	480.0	1.000000	0.000000	1.0
1.00				
Securities Account	480.0	0.125000	0.331064	0.0

```

0.00
CD Account      480.0      0.291667      0.455004      0.0
0.00
Online          480.0      0.606250      0.489090      0.0
0.00
CreditCard      480.0      0.297917      0.457820      0.0
0.00

```

```

          50%      75%      max
ID      2342.0    3566.0000  4981.0
Age      45.0     55.0000   65.0
Experience 20.0     30.0000   41.0
Income   142.5    172.0000  203.0
ZIP Code 93407.0    94705.5000 96008.0
Family   3.0      4.0000    4.0
CCAvg    3.8      5.3475    10.0
Education 2.0      3.0000    3.0
Mortgage 0.0     192.5000  617.0
Personal Loan 1.0     1.0000    1.0
Securities Account 0.0     0.0000    1.0
CD Account 0.0     1.0000    1.0
Online    1.0     1.0000    1.0
CreditCard 0.0     1.0000    1.0

```

```
no_personal_loans.describe().T
```

```


          count      mean      std      min
25% \
ID      4520.0    2512.165487  1448.299331   1.0
1259.75
Age      4520.0     45.367257   11.450427   23.0
35.00
Experience 4520.0     20.132301   11.456672  -3.0
10.00
Income   4520.0     66.237389   40.578534    8.0
35.00
ZIP Code 4520.0   93152.428761  2156.949654  9307.0
91911.00
Family   4520.0      2.373451    1.148771    1.0
1.00
CCAvg    4520.0      1.729009    1.567647    0.0
0.60
Education 4520.0      1.843584    0.839975    1.0
1.00
Mortgage 4520.0     51.789381   92.038931    0.0
0.00
Personal Loan 4520.0     0.000000    0.000000    0.0
0.00
Securities Account 4520.0     0.102212    0.302961    0.0
0.00

```


CD Account	4520.0	0.035841	0.185913	0.0
0.00				
Online	4520.0	0.595796	0.490792	0.0
0.00				
CreditCard	4520.0	0.293584	0.455454	0.0
0.00				


	50%	75%	max
ID	2518.5	3768.25	5000.0
Age	45.0	55.00	67.0
Experience	20.0	30.00	43.0
Income	59.0	84.00	224.0
ZIP Code	93437.0	94608.00	96651.0
Family	2.0	3.00	4.0
CCAvg	1.4	2.30	8.8
Education	2.0	3.00	3.0
Mortgage	0.0	98.00	635.0
Personal Loan	0.0	0.00	0.0
Securities Account	0.0	0.00	1.0
CD Account	0.0	0.00	1.0
Online	1.0	1.00	1.0
CreditCard	0.0	1.00	1.0

```
plt.figure(figsize=(15,8))
sns.distplot(personal_loans["Income"], label='Approved')
sns.distplot(no_personal_loans["Income"], label='Not Approved')
plt.legend()
plt.savefig('approved_not_approved.png', facecolor='w',
bbox_inches='tight')
plt.show()
```




```
cm = bank_df.corr()  
plt.figure(figsize=(20,20))  
sns.heatmap(cm, annot=True)  
plt.savefig('heatmap.png', facecolor='w', bbox_inches='tight')  
plt.show()
```

```
plt.figure(figsize=(15,8))
sns.distplot(bank_df["CCAvg"])
plt.show()
```



```
plt.figure(figsize=(15,8))
sns.distplot(personal_loans["CCAvg"])
sns.distplot(no_personal_loans["CCAvg"])
plt.show()
```



Data Preparation

```
from tensorflow.keras.utils import to_categorical

X = bank_df.drop(columns=["Personal Loan"])
y = bank_df["Personal Loan"]

y = to_categorical(y)

from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.1)

sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

X_train.shape, X_test.shape, y_train.shape, y_test.shape

((4500, 13), (500, 13), (4500, 2), (500, 2))
```

Building a multi-layer neural network model

```
# sequential model
ann_model = keras.Sequential()

# adding dense layer
ann_model.add(Dense(250, input_dim=13, kernel_initializer='normal',
```

```

activation='relu'))
ann_model.add(Dropout(0.3))
ann_model.add(Dense(500, activation='relu'))
ann_model.add(Dropout(0.3))
ann_model.add(Dense(500, activation='relu'))
ann_model.add(Dropout(0.3))
ann_model.add(Dense(500, activation='relu'))
ann_model.add(Dropout(0.4))
ann_model.add(Dense(250, activation='linear'))
ann_model.add(Dropout(0.4))

# adding dense layer with softmax activation/output layer
ann_model.add(Dense(2, activation='softmax'))
ann_model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
dense_6 (Dense)	(None, 250)	3500
dropout_5 (Dropout)	(None, 250)	0
dense_7 (Dense)	(None, 500)	125500
dropout_6 (Dropout)	(None, 500)	0
dense_8 (Dense)	(None, 500)	250500
dropout_7 (Dropout)	(None, 500)	0
dense_9 (Dense)	(None, 500)	250500
dropout_8 (Dropout)	(None, 500)	0
dense_10 (Dense)	(None, 250)	125250
dropout_9 (Dropout)	(None, 250)	0
dense_11 (Dense)	(None, 2)	502
=====		
Total params: 755,752		
Trainable params: 755,752		
Non-trainable params: 0		

Compilation and training of deep learning model

```
from keras import backend as K

def recall_m(y_true, y_pred):
    true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
    possible_positives = K.sum(K.round(K.clip(y_true, 0, 1)))
    recall = true_positives / (possible_positives + K.epsilon())
    return recall

def precision_m(y_true, y_pred):
    true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
    predicted_positives = K.sum(K.round(K.clip(y_pred, 0, 1)))
    precision = true_positives / (predicted_positives + K.epsilon())
    return precision

def f1_m(y_true, y_pred):
    precision = precision_m(y_true, y_pred)
    recall = recall_m(y_true, y_pred)
    return 2*((precision*recall)/(precision+recall+K.epsilon()))

ann_model.compile(loss='categorical_crossentropy', optimizer='adam',
metrics=[f1_m]) # metrics=['accuracy']

history = ann_model.fit(X_train, y_train, epochs=20,
validation_split=0.2, verbose=1)

Epoch 1/20
113/113 [=====] - 3s 14ms/step - loss: 0.2799
- f1_m: 0.9006 - val_loss: 0.1060 - val_f1_m: 0.9537
Epoch 2/20
113/113 [=====] - 1s 11ms/step - loss: 0.0927
- f1_m: 0.9710 - val_loss: 0.0887 - val_f1_m: 0.9655
Epoch 3/20
113/113 [=====] - 1s 11ms/step - loss: 0.0805
- f1_m: 0.9745 - val_loss: 0.0727 - val_f1_m: 0.9688
Epoch 4/20
113/113 [=====] - 1s 11ms/step - loss: 0.0689
- f1_m: 0.9781 - val_loss: 0.0824 - val_f1_m: 0.9666
Epoch 5/20
113/113 [=====] - 1s 11ms/step - loss: 0.0637
- f1_m: 0.9758 - val_loss: 0.0732 - val_f1_m: 0.9677
Epoch 6/20
113/113 [=====] - 1s 11ms/step - loss: 0.0649
- f1_m: 0.9754 - val_loss: 0.0767 - val_f1_m: 0.9709
Epoch 7/20
113/113 [=====] - 1s 10ms/step - loss: 0.0674
- f1_m: 0.9785 - val_loss: 0.0690 - val_f1_m: 0.9741
Epoch 8/20
```

```

113/113 [=====] - 1s 12ms/step - loss: 0.0535
- f1_m: 0.9840 - val_loss: 0.0619 - val_f1_m: 0.9784
Epoch 9/20
113/113 [=====] - 1s 11ms/step - loss: 0.0426
- f1_m: 0.9823 - val_loss: 0.0718 - val_f1_m: 0.9741
Epoch 10/20
113/113 [=====] - 1s 11ms/step - loss: 0.0615
- f1_m: 0.9804 - val_loss: 0.0742 - val_f1_m: 0.9698
Epoch 11/20
113/113 [=====] - 1s 12ms/step - loss: 0.0487
- f1_m: 0.9811 - val_loss: 0.0781 - val_f1_m: 0.9752
Epoch 12/20
113/113 [=====] - 1s 12ms/step - loss: 0.0436
- f1_m: 0.9818 - val_loss: 0.0606 - val_f1_m: 0.9795
Epoch 13/20
113/113 [=====] - 1s 12ms/step - loss: 0.0374
- f1_m: 0.9905 - val_loss: 0.0543 - val_f1_m: 0.9828
Epoch 14/20
113/113 [=====] - 1s 12ms/step - loss: 0.0395
- f1_m: 0.9845 - val_loss: 0.0602 - val_f1_m: 0.9784
Epoch 15/20
113/113 [=====] - 1s 12ms/step - loss: 0.0392
- f1_m: 0.9835 - val_loss: 0.0647 - val_f1_m: 0.9752
Epoch 16/20
113/113 [=====] - 1s 11ms/step - loss: 0.0365
- f1_m: 0.9863 - val_loss: 0.0713 - val_f1_m: 0.9795
Epoch 17/20
113/113 [=====] - 1s 11ms/step - loss: 0.0328
- f1_m: 0.9895 - val_loss: 0.0605 - val_f1_m: 0.9752
Epoch 18/20
113/113 [=====] - 1s 11ms/step - loss: 0.0386
- f1_m: 0.9832 - val_loss: 0.0766 - val_f1_m: 0.9784
Epoch 19/20
113/113 [=====] - 1s 11ms/step - loss: 0.0340
- f1_m: 0.9901 - val_loss: 0.0743 - val_f1_m: 0.9774
Epoch 20/20
113/113 [=====] - 1s 11ms/step - loss: 0.0383
- f1_m: 0.9851 - val_loss: 0.1021 - val_f1_m: 0.9731

```

```


# Plot the model performance across epochs

```

```

plt.figure(figsize=(15,8))
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('loss')
plt.legend(['train_loss', 'val_loss'], loc = 'upper right')
plt.savefig('modelloss.png', facecolor='w', bbox_inches='tight')
plt.show()

```

Evaluating model performance

```
predictions = ann_model.predict(X_test)
predict = []

for i in predictions:
    predict.append(np.argmax(i))

from sklearn import metrics
y_test = np.argmax(y_test, axis=1)


f1_test = metrics.f1_score(y_test, predict)
prec = metrics.precision_score(y_test, predict)
rec = metrics.recall_score(y_test, predict)
acc = metrics.accuracy_score(y_test, predict)

print ("F1 Score: {:.4f}.".format(f1_test))
print ("Precision: {:.4f}.".format(prec))
print ("Recall: {:.4f}.".format(rec))
print ("Accuracy: {:.4f}.".format(acc)) # note this is not a good
measure of performance for this project as dataset is unbalanced.

F1 Score: 0.8539.
Precision: 0.9500.
Recall: 0.7755.
Accuracy: 0.9740.

conf_mat = metrics.confusion_matrix(y_test, predict)
plt.figure(figsize=(10,8))
sns.heatmap(conf_mat, annot=True, cbar=False)
```

```
plt.savefig('conf_matrix.png', facecolor='w', bbox_inches='tight')
plt.show()
```



```
print(metrics.classification_report(y_test, predict))
```

	precision	recall	f1-score	support
0	0.98	1.00	0.99	451
1	0.95	0.78	0.85	49
accuracy			0.97	500
macro avg	0.96	0.89	0.92	500
weighted avg	0.97	0.97	0.97	500